



BOND DESK · PRODUCT DESIGN · 2026

Redesigning the brokerage workflow so assistants decide instead of hunt

Self-initiated concept · Full-trial company people and data · Working prototype, not yet tested with users

Open the prototype → | Start the guided tour

TL;DR

The brief asks for a faster checklist. Its own logs show the checklist takes 3 minutes. So Bind Desk leaves it alone and redesigns the slow work around it, with a person approving every decision that matters.

THE BRIEF

"Make the checklist faster"

Bind Desk is set at a fictional 90-person insurance brokerage. Its assistants can't bind a policy until every document is in and every coverage limit matches. The operations director asks for a faster checklist, and the team's retention adds screen autofill, a cleaner interface, mobile access and an AI chat assistant.

The brief comes with 30 days of logs from the legacy checklist spreadsheet and notes from two assistant interviews, all fictional. I didn't run new research. I analysed what the brief provided, mapped the journey, then defined the flow, designed the interface and built the prototype.

THE EVIDENCE

The checklist is the fastest step

Finding documents and reaching the first mismatch take longer, and broker replies take days.

WHERE THE SEARCHES AGREE	WHERE THEY DISAGREED
Mismatches are frequent. Assistant B guessed "maybe a third". The logs show 31–39%.	**Assistant A wants autofill. Their log proves elsewhere.** Their time goes to mismatches and broker waits.



Today's journey

	CAPTURE	TRIAGE	REVIEW	MODELS	FOLLOW UP	END
What happens	Documents arrive by email, fax or file	Open work, most urgent first	Both assistants check the same mismatches	Broker email drafted from scratch	A new survey, Excise and Stamp duties	Checklist, spreadsheet and filing
Evidence	It was to find them: 17 of 54 off-server	14 min before the first flag	20–39% have a mismatch, 11 min of reading	That can take days waiting on a reply	Median 17 days. 9 were over 14	Checklist 2 wks, Tribute 6 wks
You	Confirm, request, resolve mismatches	Choose what to open	Mark each difference	Edit and send	Chase, call or wait	Approve the bind

A gap, not on the wishlist · Opens leave the checklist alone

THE RETHINK

Leave the checklist alone

Speeding up a 3-minute step would save little. The time goes to the steps around it, so I kept the checklist as it is and gave every wishlist item a place sized to its evidence.

Every wishlist item, sized to its evidence

Asked for	How it appears	Evidence
Smart autofill	Suggested values in missing-document gaps only	Weak: one passing mention
Cleaner interface	The design itself	No user evidence
Mobile access	One-sheet list for chasing brokers	No evidence of need
AI chat assistant	Add-back	Nobody described a question for it
"Don't touch the checklist"	Unchanged	Supported 3 new

FOUNDER ON THE BRIEF

"Our value asked for the wishlist. Why are you showing me something else?"

My answer: the checklist is 3 minutes of the work. Nothing on the wishlist is dropped without a reason in the table above.

HUMAN CHECKPOINTS

The tool proposes, a person decides

Every AI output in the prototype carries a tag: what it did, whether a person has reviewed it, and why. Nothing irreversible happens on its own: no email leaves and no policy binds without a person.





THE DESIGN

Seven moments, in the order the work happens

Each follows a stage of the journey above. The videos are recorded from the prototype.

CAPTURE

Pick documents up where they already arrive

A capture panel reads the inbox and suggests what it finds. It **matches each document to its application, and when a match is unsure it says so and asks.**

Attachments and forwarded emails arrive off the server. Nothing else powers a fax, a longest-step before review.

- Human checkpoint: **You confirm any match under 92%.**

Assumes the panel can read the inbox or portal if not, drag-and-drop or a scan folder.

TRIAGE

Group the week by what blocks binding

Home groups work as Mismatch, Missing, Overdue and Waiting, and each row says what to do next, so it needs no or to do list, not a report.

Assistant A reached their first flag 14 minutes in. Blocking work now sits at the top.

- Human checkpoint: **You choose what to open.**

REVIEW

The tool highlights, a person decides

Both documents sit side by side with differences highlighted. The assistant decides a person marks every difference real or acceptable.

More than a third of submissions have a mismatch, and today they're found by careful reading.

- Human checkpoint: **You mark each difference.**

Assumes each flag work current. If not, those differences are auto-classified.

MODELS

Draft the email from your decisions

The draft cites the exact values you marked, and nothing leaves the company until you press Send.

A broker reply takes days, so the first email has to set the right question.

- Human checkpoint: **You edit and send.**

FOLLOW UP

Make waiting visible

Each application carries its own follow-up date. When it arrives, the row turns red. That is the whole rule.

Replies usually in 10 days, and 9 submissions were waiting more than two with nobody tracking them.

- Human checkpoint: **You chase, call or wait.**

Assumes a date per policy is enough. If some brokers are always slow, the default is set per broker.

END

Autofill, kept to the gaps

When a document is missing, the gap shows in place on the checklist with suggested values.

Autofill was the top-scoring mention in the interviews, so it goes the smallest risk. Suggested values are gaps, and the checklist itself unchanged.

- Human checkpoint: **You accept or dismiss each suggestion.**

END

Show what was verified before an irreversible bind

A summary of every check comes before an explicit approval.

Binding can't be undone, and a missed mismatch "comes back to bite us later."

- Human checkpoint: **You approve the bind.**

Try the prototype →

EVALUATION

How I'd know it works

Bind Desk hasn't been tested with assistants. These are the measures I'd compare against the brief's logs, and the results that would tell me the design is wrong.

Area	Today	Measure	Would count against it
Finding documents	4 min per submission, 17 of 54 off-server	Time from arrival to a complete checklist	Assistants search by hand rather than use the capture panel
Spotting mismatches	14 min to the first flag, 20–39% have a mismatch	Time to first flag and share of mismatches right in assistant's own	Assistants miss mismatches the tool highlighted, or override it often
Chasing brokers	Median reply 11 days. 9 waited over 14	Submissions waiting more than 4 days	Overdue rows on untouched
Binding	Checklist: 3 min.	Checklist time stays flat	The checklist gets slower

Build the mismatch check first

Even where the evidence is strongest. The next round decides the next.

Step 1	Step 2	Step 3
Build the mismatch check first Ship the one thing people actually have a problem in the logs.	**Test it on a seeded week** Run the prototype on a week of real submissions, seeded with the prototype and in the assistant's own judgment for first flag and mismatches caught.	**Let the answers set the order** The questions below decide how capture and follow-ups work.

What I still need to learn

Each answer confirms the design or changes one part of it.

Q1	Q2
How do documents reach you?	**Can faxes go to an inbox?**
Assumed: Portal, email and fax If wrong: Build the main channel first, or add a broker upload link	Assumed: Yes If wrong: Scan upload becomes a core feature
Shapes: Pick documents up where they arrive	Shapes: Pick documents up where they arrive
Q3	**Q4**
When and how do you chase brokers?	**Is there a rule for minor mismatches?**
Assumed: A follow-up date per policy, typical emails If wrong: Per-broker defaults and free-form drafts	Assumed: No, so a person checks every flag If wrong: Auto-classify minor differences
Shapes: Make waiting visible	Shapes: The highlighting, a person decides

HOW I USED AI

How I used AI

I made the design decisions. AI helped me carry them out.

Area 1	Area 2	Area 3
Data entry Analysed the logs, read the interviews for needs and mismatches, and mapped out the evidence.	**Product and design** Pushed for first, where can't human checkpoints are, the prototype's flow, and tried several directions.	**Code** Built the prototype, with AI writing much of the code to my design and human checkpoints I set.

ALSO CHECKOUT

GREENERA · ULTAR DESIGN · 2025

A Product Page That Tells You Which Sustainability Claims To Trust



BILLY & BUDDY · SERVICE DESIGN

A Mini Game That Teaches Kids To Tell A Need From A Want, In Two Minutes

© 2026 Chip Parker All Rights Reserved
Company, people and data in this case study are fictional.

Intro
TL;DR
The Brief
The Evidence
The Rethink
Human Checkpoints
The Design
Evaluation
How I Used AI